

3-ways to Set up LLaMA 2 Locally on CPU (Part 3 — Hugging Face)



Antoine Frd · [Follow](#)

5 min read · Jan 31, 2024



Listen



Share

... More



- [Load LLaMA 2 model with Hugging Face](#)
- [Install dependencies for running Llama 2 with Hugging Face locally](#)
- [Downloading Llama 2 model](#)
- [Running the model using Hugging Face pipeline](#)

Load LLaMA 2 model with Hugging Face

Install dependencies for running Llama 2 with Hugging Face locally

We need to ensure that the essential libraries are installed:

- `transformers` : Hugging Face Transformers provides APIs to quickly download and use pretrained Transformer models.
- `torch` : PyTorch provides Tensors that can live either on the CPU or the GPU

Note: PyTorch is a **GPU-Ready Tensor library**. For CPU-support only, we can conveniently install Hugging Face Transformers and PyTorch library in one command line

```
pip install 'transformers[torch]'
```

Explanation:

- `pip install torch transformers` would install two packages named `torch` and `transformers`.
- `pip install transformers[torch]`, on the other hand, installs a variant of the `transformers` package which contains support for `torch`. Note that it has nothing to do with the `torch` package itself, but is just a string defined by the `transformers` package for a particular **feature set** that gets enabled. How the argument `torch` is interpreted is entirely up to its `setup.py` file. (Brackets [optional] in PIP).

1° First, Download the library.

```
pip3 install 'transformers[torch]'
```

2° Create an Account on Hugging Face

- Open a new tab and visit the Hugging Face website.
- Click on the “**Sign In**” button or sign up to Hugging Face.
- Once you have created your Hugging Face account, log in to the platform.



Join Hugging Face

Join the community of machine learners!

Email Address

Hint: Use your organization email to easily find and join your company/team org.

Password

Next

Already have an account? [Log in](#)

[SSO](#) is available for companies

Obtaining the Hugging Face API Token

- Click on your profile picture at the top right corner of the page.
- Select “**Access Token**” from the dropdown menu.

Profile

Account

Organizations

Billing

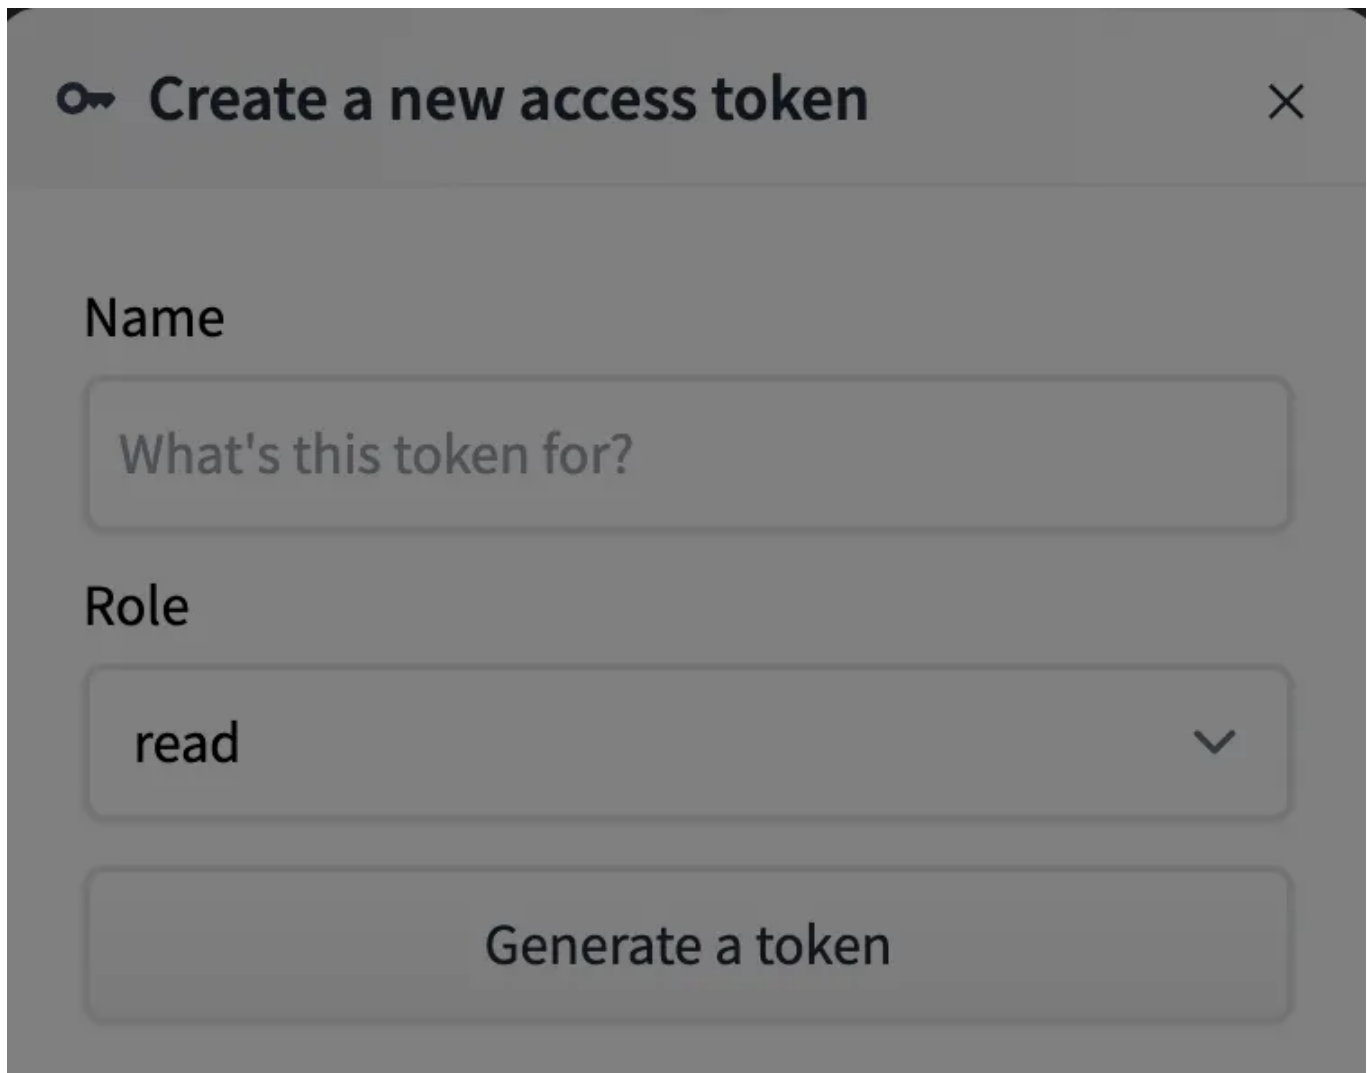
Access Tokens

SSH and GPG Keys

Webhooks

Papers

- Click on the “**New Token**” button.
- Give your token a name and click on the “**Generate a token**” button.



Create a new access token

Name

What's this token for?

Role

read

Generate a token

- Copy the Hugging Face API token. This token will be used to authenticate your requests to the Hugging Face API.

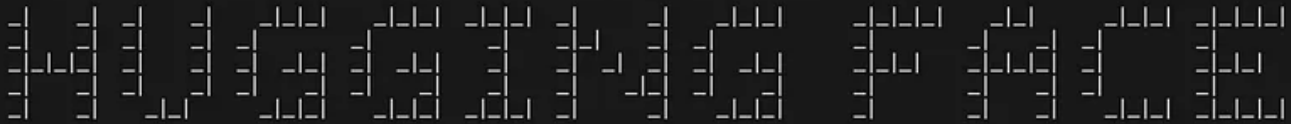
3° Login

Once you have an account and your token, run the following command in your terminal or in your notebook:

```
huggingface-cli login  
  
#or  
  
huggingface-cli login --token YOUR_TOKEN_HERE
```

This command will prompt you to enter your Hugging Face token.

```
(.venv) antoine@antoinet-mbp Install-Llama2-locally % huggingface-cli login
```



To login, `huggingface_hub` requires a token generated from <https://huggingface.co/settings/tokens> .
Token:

You can use the following command to check if you are logged and determine your currently active account

```
huggingface-cli whoami
```

Note: Alternatively, you can login using `login()` or `login_notebook()` in a script or a notebook:

```
#!/pip3 install --upgrade huggingface_hub

from huggingface_hub import notebook_login
notebook_login()
```

You can only be logged in to one account at a time. (see [more](#))



Copy a token from [your Hugging Face tokens page](#) and paste it below.

Immediately click login after copying your token or it might be stored in plain text in this notebook file.

Token:

☒ Add token as git credential?

Login

Pro Tip: If you don't already have one, you can create a dedicated 'notebooks' token with 'write' access, that you can then easily reuse for all notebooks.

Downloading Llama 2 model

Note: Compared with the model used in the first part [llama-2-7b-chat.Q2_K.gguf](#) (Part. 1). We cannot use the `transformers` library. GGML and GGUF models are **not natively supported** in `transformers` core. Instead of importing `AutoModel` from `transformers` we can import it from the `ctransformers` library to deal with GGML or GGUF models.

We gonna use the [meta-llama/Llama-2-7b-hf](#) model.

Note: In order to use Llama-2 with Hugging Face, you need to raise a [request](#) on the model page. (Make sure you are using the same email ids in both places).



Hi Frd,

This is to let you know your request to access model "meta-llama/Llama-2-70b-chat" on [huggingface.co](#) has been accepted by the repo authors.

You can now access its content [here](#).

Cheers,

The Hugging Face team

There are several ways to download the model from Hugging Face to use it locally. Here are 3 ways to do it:

Method 1: Use `from_pretrained()` and `save_pretrained()` HF functions

- Download required files:

```
from transformers import AutoTokenizer, AutoModel
```

```
#Download your files online
```

```
tokenizer = AutoTokenizer.from_pretrained("meta-llama/Llama-2-7b-chat-hf")
```

```
model = AutoModel.from_pretrained("meta-llama/Llama-2-7b-chat-hf")
```

- Use the `save_pretrained()` function to download a file to a specific local path.

```
#Save your files to a specified directory
```

```
tokenizer.save_pretrained("./models/hf-frompretrained-download/Llama-2-7b-chat-hf/")
```

```
model.save_pretrained("./models/hf-frompretrained-download/Llama-2-7b-chat-hf/")
```

- This method allow us to download only the required files from the model repository.

✓ models

✓ hf-frompretrained-download / Llama-2-7b-chat-hf

{ } config.json

≡ model-00001-of-00006.safetensors

≡ model-00002-of-00006.safetensors

≡ model-00003-of-00006.safetensors

≡ model-00004-of-00006.safetensors

≡ model-00005-of-00006.safetensors

≡ model-00006-of-00006.safetensors

{ } model.safetensors.index.json

{ } special_tokens_map.json

{ } tokenizer_config.json

{ } tokenizer.json

≡ tokenizer.model

- Then, use the `from_pretrained()` function to download the model offline.


```
#Reload your files offline
tokenizer = AutoTokenizer.from_pretrained("./models/hf-frompretrained-download/Llama-2-7b-chat-hf")
model = AutoModel.from_pretrained("./models/hf-frompretrained-download/Llama-2-7b-chat-hf")
```

Method 2: Use `snapshot_download()` HF functions

- Download a snapshot of the model repository:

```
MODEL_REPO = "meta-llama/Llama-2-7b-chat-hf"

model_path = snapshot_download(
    repo_id=MODEL_REPO,
    local_dir="./models/hf-snapshot-download/Llama-2-7b-chat-hf",
    cache_dir="./models/hf-snapshot-download/Llama-2-7b-chat-hf",
    local_files_only = False,
)
```

Unlike the previous method, this one download the entire content of the model's repository. We can filter the files extensions using ``allow_patterns`` and ``ignore_patterns`` parameters.

```
✓ hf-snapshot-download / Llama-2-7b-chat-hf
> .locks
> models--meta-llama--Llama-2-7b-chat-hf
🔗 .gitattributes
{} config.json
{} generation_config.json
🔑 LICENSE.txt
≡ model-00001-of-00002.safetensors
≡ model-00002-of-00002.safetensors
{} model.safetensors.index.json
≡ pytorch_model-00001-of-00002.bin
≡ pytorch_model-00002-of-00002.bin
{} pytorch_model.bin.index.json
📘 README.md
{} special_tokens_map.json
{} tokenizer_config.json
{} tokenizer.json
≡ tokenizer.model
📄 USE_POLICY.md
```

Method 3: Use `git clone`

- Clone the repository

```
cd models
cd hf-git-clone
git init
git lfs install
git clone <https://huggingface.co/meta-llama/Llama-2-7b-chat-hf>
```

Note: If Git credential manager asks you for your authentication credentials, use your token as `https://<user_name>:<token>@huggingface.co/<repo_path>` .(see [more](#))

Running the model using Hugging Face pipeline

Once the model downloaded, we can generate responses using the Hugging Face Transformers Pipelines

```
from transformers import pipeline
import torch

model = "./models/hf-snapshot-download/Llama-2-7b-chat-hf"

llama_pipeline = pipeline(
    "text-generation", # LLM task
    model=model,
    torch_dtype=torch.bfloat16,
    device = -1, # or "cpu"
)
```

```
def get_llama_response(prompt: str) -> None:
    """
    Generate a response from the Llama model.

    Parameters:
        prompt (str): The user's input/question for the model.

    Returns:
        None: Prints the model's response.
    """

    sequences = llama_pipeline(
        prompt,
        do_sample=True,
        top_k=10,
        num_return_sequences=1,
        eos_token_id=tokenizer.eos_token_id,
        max_length=256,
    )

    print("Llama 2 Chatbot:", sequences[0]['generated_text'])
```

Open in app ↗



Search



Llama 2 Chatbot: Which planet is the closest to the sun?

Answer: Mercury is the closest planet to the sun, with an average distance of about 58 mill-

Congratulations! In this third and last article we've successfully installed Llama 2 with Hugging Face. We can now use it locally with CPU only.

[Llm](#)[Generative Ai Use Cases](#)[Hugging Face](#)[Llama 2](#)[Cpu](#)[Follow](#)

Written by Antoine Frd

16 Followers

AI Engineer

More from Antoine Frd