

了解 LLaMA-2 模型结构 (5)

发表评论 / ChatGPT, GPT, OpenAI

9. 转换 tokenizer.model 并保存

前面的章节转换模型的所有权重后，还需要转换 tokenizer.model 为自己需要的格式。

把 meta-llama/Llama-2-7b-chat-hf/ 目录下的 tokenizer.model 拷贝到 newsrc 目录下。

参照 <https://github.com/karpathy/llama2.c> 项目下的 tokenizer.py 文件，命名为 test09.py，文件保存到 newsrc 目录下：

```
1 import os
2 import struct
3 import argparse
4 from typing import List
5
6 from sentencepiece import SentencePieceProcessor
7
8 class Tokenizer:
9     def __init__(self, tokenizer_model=None):
10         model_path = tokenizer_model
11         assert os.path.isfile(model_path), model_path
12         self.sp_model = SentencePieceProcessor(model_file=model_path)
13         self.model_path = model_path
14
15         # BOS / EOS token IDs
16         self.n_words: int = self.sp_model.vocab_size()
17         self.bos_id: int = self.sp_model.bos_id()
18         self.eos_id: int = self.sp_model.eos_id()
19         self.pad_id: int = self.sp_model.pad_id()
20         #print(f"#words: {self.n_words} - BOS ID: {self.bos_id} - EOS ID: {self.eos_id}")
21         assert self.sp_model.vocab_size() == self.sp_model.get_piece_size()
22
23     def encode(self, s: str, bos: bool, eos: bool) -> List[int]:
24         assert type(s) is str
25         t = self.sp_model.encode(s)
26         if bos:
27             t = [self.bos_id] + t
28         if eos:
29             t = t + [self.eos_id]
30         return t
31
32     def decode(self, t: List[int]) -> str:
33         return self.sp_model.decode(t)
34
35     def export(self):
36
37         # get all the tokens (postprocessed) and their scores as floats
38         tokens, scores = [], []
39         for i in range(self.n_words):
40
41             # decode the token and light postprocessing
42             t = self.sp_model.id_to_piece(i)
43             s = self.sp_model.get_score(i)
44             if i == self.bos_id:
45                 t = '\n<s>\n'
46             elif i == self.eos_id:
```

```

47         t = '\n</s>\n'
48         t = t.replace('_', ' ') # sentencepiece uses this character as whitespace
49         b = t.encode('utf-8') # bytes of this token, utf-8 encoded
50
51         tokens.append(b)
52         scores.append(s)
53
54         # record the max token length
55         max_token_length = max(len(t) for t in tokens)
56
57         # write to a binary file
58         # the tokenizer.bin file is the same as .model file, but .bin
59         tokenizer_bin = self.model_path.replace('.model', '.bin')
60         with open(tokenizer_bin, 'wb') as f:
61             f.write(struct.pack("I", max_token_length))
62             for bytes, score in zip(tokens, scores):
63                 f.write(struct.pack("fI", score, len(bytes)))
64             f.write(bytes)
65
66 t = Tokenizer("newsrsrc/tokenizer.model")
67 t.export()

```

运行 test09.py, 查看newsrsrc 的文件目录

```

1 ls -l newsrsrc/tokenizer.*
2 -rwxrwxrwx 1 tony tony 433869 Mar 11 14:24 newsrsrc/tokenizer.bin
3 -rwxrwxrwx 1 tony tony 499723 Mar 10 23:51 newsrsrc/tokenizer.model

```

10. 查看 tokenizer.bin

下面给出的例子开始的都是C/C++代码, 这样更好理解文件里面的内容

参照 <https://github.com/karpathy/llama2.c> 项目下的 run.c 文件, 命名为 test01.c, 文件保存到 newsrsrc 目录下:

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 // The Byte Pair Encoding (BPE) Tokenizer that translates strings <-> tokens
4
5 typedef struct {
6     // 一个字符指针, 指向与ID关联的字符串。
7     char *str;
8     // 一个整数, 表示与字符串关联的ID。
9     int id;
10 } TokenIndex;
11
12 typedef struct {
13     // 一个指针数组, 存储词汇表中每个单词的字符串表示。
14     // 例如, 如果 vocab[0] 是 "apple", 那么 vocab[0][0] 就是字符 'a'。
15     char** vocab;
16     // 一个浮点数数组, 可能存储与词汇表中的每个单词相关联的分数或权重。
17     float* vocab_scores;
18     // 结构体, 存储一个字符串和与其相关联的整数ID
19     TokenIndex *sorted_vocab;
20     // 一个整数, 表示词汇表的大小, 即vocab和vocab_scores数组中的元素数量。
21     int vocab_size;
22     // 一个无符号整数, 表示词汇表中的最大token长度。
23     unsigned int max_token_length;
24     // 一个无符号字符数组, 存储所有单字节字符串
25     // 这个数组的大小被设定为512, 可能是为了存储ASCII字符表中的所有可能的单字节字符串。
26     unsigned char byte_pieces[512]; // stores all single-byte strings
27 } Tokenizer;

```

```

28
29 void build_tokenizer(Tokenizer* t, char* tokenizer_path, int vocab_size) {
30     // i should have written the vocab_size into the tokenizer file... sigh
31     // 设置了Tokenizer结构体中vocab_size的值。
32     t->vocab_size = vocab_size;
33     // malloc space to hold the scores and the strings
34     // 为vocab, vocab_scores数组分配了内存, 而sorted_vocab被初始化为NULL, 表示它将在后续被“懒惰地”初始化
35     t->vocab = (char**)malloc(vocab_size * sizeof(char*));
36     t->vocab_scores = (float*)malloc(vocab_size * sizeof(float));
37     t->sorted_vocab = NULL; // initialized lazily
38     // 初始化byte_pieces数组, 该数组存储所有单字节字符串
39     for (int i = 0; i < 256; i++) {
40         t->byte_pieces[i * 2] = (unsigned char)i;
41         t->byte_pieces[i * 2 + 1] = '\0';
42     }
43     // read in the file
44     FILE *file = fopen(tokenizer_path, "rb");
45     if (!file) { fprintf(stderr, "couldn't load %s\n", tokenizer_path); exit(EXIT_FAILURE); }
46     // 从文件中读取max_token_length的值
47     if (fread(&t->max_token_length, sizeof(int), 1, file) != 1) { fprintf(stderr, "failed read\n"); exit(EXIT_FAILURE); }
48     int len;
49     // 这个循环块读取每个vocab字符串及其对应的vocab_scores值。
50     for (int i = 0; i < vocab_size; i++) {
51         if (fread(t->vocab_scores + i, sizeof(float), 1, file) != 1) { fprintf(stderr, "failed read\n"); exit(EXIT_FAILURE); }
52         if (fread(&len, sizeof(int), 1, file) != 1) { fprintf(stderr, "failed read\n"); exit(EXIT_FAILURE); }
53         t->vocab[i] = (char *)malloc(len + 1);
54         if (fread(t->vocab[i], len, 1, file) != 1) { fprintf(stderr, "failed read\n"); exit(EXIT_FAILURE); }
55         t->vocab[i][len] = '\0'; // add the string terminating token
56     }
57     fclose(file);
58 }
59
60 void free_tokenizer(Tokenizer* t) {
61     for (int i = 0; i < t->vocab_size; i++)
62     {
63         free(t->vocab[i]);
64     }
65     if(t->vocab) free(t->vocab);
66     if(t->vocab_scores) free(t->vocab_scores);
67     if(t->sorted_vocab) free(t->sorted_vocab);
68 }
69
70 void print_tokenizer(Tokenizer* t) {
71     printf("vocab = %d\n", t->vocab_size);
72     printf("max_token_length = %d\n", t->max_token_length);
73     for (int i = 0; i < t->vocab_size; i++)
74     {
75         printf("%5d, %12.6lf, (%s)\n", i, t->vocab_scores[i], t->vocab[i]);
76     }
77 }
78
79 int main(int argc, char *argv[]) {
80
81     char *tokenizer_path = "tokenizer.bin";
82
83     // build the Tokenizer via the tokenizer .bin file
84     Tokenizer tokenizer;
85     int vocab_size = 32000; // 从模型文件的 config.json 获取
86     build_tokenizer(&tokenizer, tokenizer_path, vocab_size);
87
88     print_tokenizer(&tokenizer);
89     free_tokenizer(&tokenizer);
90
91     return 0;
92 }

```

编译 test01.c

```
1 make test01
2 cc      test01.c  -o test01
```

运行 test01

```
1 ./test01 > 1.txt
```

由于输出的内容很多，所以我们把输出重定向到 1.txt 文件中，下面是 1.txt 文件的开头和结尾部分内容

```
1 vocab = 32000
2 max_token_length = 27
3 0, 0.000000, (<unk>)
4 1, 0.000000, (
5 <s>
6 )
7 2, 0.000000, (
8 </s>
9 )
10 3, 0.000000, (<0x00>)
11 4, 0.000000, (<0x01>)
12 5, 0.000000, (<0x02>)
13 6, 0.000000, (<0x03>)
14 7, 0.000000, (<0x04>)
15 8, 0.000000, (<0x05>)
16 9, 0.000000, (<0x06>)
17 10, 0.000000, (<0x07>)
18 11, 0.000000, (<0x08>)
19 12, 0.000000, (<0x09>)
20 ...
21 259, -1000000000.000000, ( )
22 260, -1.000000, ( t)
23 261, -2.000000, (er)
24 262, -3.000000, (in)
25 263, -4.000000, ( a)
26 264, -5.000000, (en)
27 265, -6.000000, (on)
28 266, -7.000000, ( th)
29 267, -8.000000, (es)
30 268, -1000000000.000000, ( )
31 269, -10.000000, ( s)
32 270, -11.000000, ( d)
33 271, -12.000000, (at)
34 ...
35 31985, -31726.000000, (怪)
36 31986, -31727.000000, (联)
37 31987, -31728.000000, (역)
38 31988, -31729.000000, (泰)
39 31989, -31730.000000, (백)
40 31990, -31731.000000, (ò)
41 31991, -31732.000000, (げ)
42 31992, -31733.000000, (ベ)
43 31993, -31734.000000, (边)
44 31994, -31735.000000, (还)
45 31995, -31736.000000, (黄)
46 31996, -31737.000000, (왕)
47 31997, -31738.000000, (收)
48 31998, -31739.000000, (弘)
49 31999, -31740.000000, (给)
```

可以看到，token 的最大长度为27